

Uni.lu HPC School 2018

PS5: Parallel computations with OpenMP/MPI



Uni.lu High Performance Computing (HPC) Team

Dr. S. Varrette

University of Luxembourg (UL), Luxembourg

<http://hpc.uni.lu>



Latest versions available on Github:



UL HPC tutorials:

<https://github.com/ULHPC/tutorials>

UL HPC School:

<http://hpc.uni.lu/hpc-school/>

PS5 tutorial sources:

ulhpc-tutorials.rtf.d.io/en/latest/parallel/basics/





Summary

- 1 Introduction**
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)

Main Objectives of this Session

- See how to run **threaded parallel OpenMP** programs
 - See how to **use the MPI suits** available on the **UL HPC** platform:
 - ↪ Intel MPI and the Intel MKL
 - ↪ OpenMPI
 - ↪ MVAPICH2
 - **Build** and **run** OpenMP and/or MPI code
 - ↪ interactive and passive (through **launcher**) jobs
- **Test cases** on reference parallel and distributed **benchmarks**:
 - ↪ OSU micro-benchmarks:
 - ✓ measure the performances of various MPI operations
 - ↪ High-Performance Linpack (HPL)
 - ↪ Hybrid High Performance Conjugate Gradients (HPCG)



Summary

- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs**
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)

OpenMP

- OpenMP: Open Multi-Processing
 - popular parallel programming model for multi-threaded applications.
 - API for **multi-platform shared memory multiprocessing**
 - ✓ in C, C++, and Fortran
 - ✓ on most platforms, instruction set architectures and OS.
 - Parallelism accomplished **exclusively** through the use of threads.
 - ✓ **Thread**: smallest unit of processing that can be scheduled by an OS
 - **#threads** \simeq **number of machine processors/cores**.
 - ✓ OMP_NUM_THREADS (if present): **initial max** number of threads;

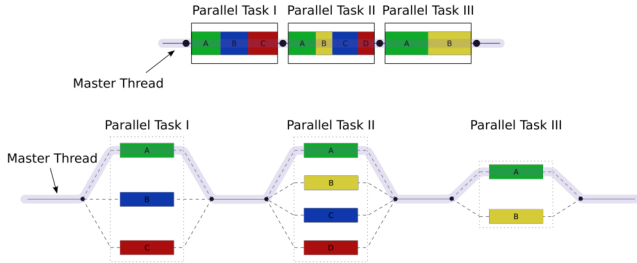
OpenMP

<https://www.openmp.org/>

- Reference website
- **Latest version: 5.0** (Nov 2018) – specifications

OpenMP Programming Model

- Explicit (not automatic) programming model
 - ↳ may mean taking a serial program & insert compiler directives. . .
- **Fork-Join** model of parallel execution
 - ↳ **FORK**: **master** thread creates a **team** of parallel threads.
 - ↳ **JOIN**: team threads complete statements in parallel regions
 - ✓ then they synchronize & terminate, leaving only the master thread.



OpenMP on UL HPC platform

- Rely on **Environment Modules** **once** on a computing node
- Comes part of the `intel` or the `foss` toolchains modules

Toolchain	Compilation command
<code>toolchain/intel</code>	<code>icc -qopenmp [...]</code>
<code>toolchain/foss</code>	<code>gcc -fopenmp [...]</code>

Reservation

- You will probably want to compile/test within an **interactive job**
↪ **Ex:** 4 threads (core) on 1 node

```
# SLURM -- Iris cluster
```

```
(access)$> srun -p interactive --ntasks-per-node=1 -c 4 --pty bash
```

```
# OAR -- gaia, chaos cluster
```

```
(access)$> oarsub -I -l nodes=1/core=4,walltime=4
```

OpenMP with OAR

- Reservation has to match the required number of OpenMP threads:

↪ **Ex:** `oarsub -l nodes=1/core=4[...]`

```
#!/bin/bash -l
#OAR -l nodes=1/core=4,walltime=1

export OMP_NUM_THREADS=$(cat $OAR_NODEFILE | wc -l)

# Use the RESIF build modules of the UL HPC platform
if [ -f /etc/profile ]; then
    . /etc/profile
fi

# Load the {intel | foss} toolchain and whatever module(s) you need
module purge
module load toolchain/intel # or toolchain/foss

path/to/your/openmp_program
```

OpenMP with Slurm

sbatch/srun option	Description
--ntasks-per-node=1	single task per node
-c <T>	number of OpenMP threads

```
#!/bin/bash -l
#SBATCH --ntasks-per-node=1      # more explicit than '-n 1'
#SBATCH -c 28                    # number of CPU cores / OpenMP threads per task
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch
if [ -f /etc/profile ]; then
    . /etc/profile
fi
# Load the {intel / foss} toolchain and whatever module(s) you need
module purge
module load toolchain/intel # or toolchain/foss
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}
srun /path/to/your/openmp_program
```

Hands-on 1: Parallel OpenMP jobs

Your Turn!

Hands-on 1

ulhpc-tutorials.rtf.d.io/en/latest/parallel/basics/#parallel-openmp-jobs

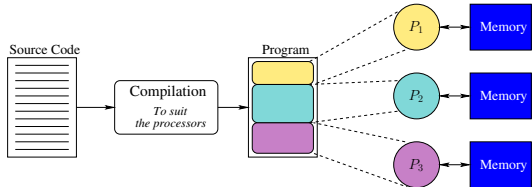
- **Reserve** an **interactive** job to launch 4 OpenMP threads
- Check and **compile** `src/hello_openmp.c`
 - ↪ against both toolchains `bin/[intel_]hello_openmp`
- **execute** the generated binaries
 - ↪ set `$OMP_NUM_THREADS`
- prepare a **launcher script** `runs/launcher.OpenMP.sh`
- repeat on a more serious program `src/matrix_mult_openmp.c`
 - ↪ `bin/[intel_]matrix_mult_openmp`
- (eventually) OpenMP **DataRaceBench** suite



Summary

- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs**
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)

SPMD Programming model



• SPMD: **Simple Program, Multiple Data**

- ↪ same programs for each processors
 - ✓ executed at independent points
- ↪ processes identified by a rank
 - ✓ each process knows the piece of code he works on
 - ✓ common in master-worker computations

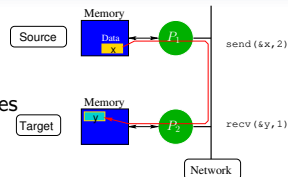
```

if (my_rank == 0) { /* master */
    // ... load input and dispatch ...
} else { /* workers */
    // ... wait for data and compute ...
}
  
```

MPI (Message Passing Interface)

- **Message Passing Model:**

- ↪ each “processor” runs a process
- ↪ processes communicate by exchanging messages
 - ✓ *analogy: mail*



Message Passing Interface (MPI) Standard

- **Goal:**

- ↪ **portable, efficient & flexible** standard for message passing
- ↪ industry standard

- **Reference website**

<https://www.mpi-forum.org/>

- **Latest version: 3.1** (June 2015) – **specifications**

MPI on UL HPC platform

- Rely on **Environment Modules** **once** on a computing node
↳ then you can search for MPI suites available: `module avail mpi`

MPI Suite	Version	module load...	Compiler
Intel MPI	18.0.1	toolchain/intel	C: mpiicc; C++: mpiicpc
OpenMPI	2.1.3	mpi/OpenMPI	C: mpicc; C++: mpic++
MVAPICH2	2.3a	mpi/MVAPICH2	C: mpicc; C++: mpic++

MPI on UL HPC platform

- Rely on **Environment Modules** **once** on a computing node
 ↳ then you can search for MPI suites available: `module avail mpi`

MPI Suite	Version	module load...	Compiler
Intel MPI	18.0.1	toolchain/intel	C: <code>mpiicc</code> ; C++: <code>mpiicpc</code>
OpenMPI	2.1.3	mpi/OpenMPI	C: <code>mpicc</code> ; C++: <code>mpic++</code>
MVAPICH2	2.3a	mpi/MVAPICH2	C: <code>mpicc</code> ; C++: <code>mpic++</code>

MPI Suite	Compilation command (in C)
Intel MPI	<code>mpiicc -Wall [-qopenmp] [-xhost] -O2 [...]</code>
OpenMPI	<code>mpicc -Wall [-fopenmp] -O2 [...]</code>
MVAPICH2	<code>mpicc -Wall [-fopenmp] -O2 [...]</code>

Reservation

- You will probably want to compile/test within an **interactive job**
↪ **Ex:** 1 core on two **different** nodes

```
# SLURM -- Iris cluster
```

```
(access)$> srun -p interactive -N 2 --ntasks-per-node 1 --pty bash
```

```
# OAR -- gaia, chaos cluster
```

```
(access)$> oarsub -I -l nodes=2/core=1,walltime=4
```

- Then you can search for MPI suites available

```
(node)$> module avail mpi
```

MPI with OAR

- Reservation has to match the required number of MPI processes:
 ↪ **Ex:** `oarsub -l nodes=2/core=12[...]`
- Running MPI code using traditional `mpirun`

```
# Intel MPI
$> module load toolchain/intel
# ONLY on moonshot node have no IB card: export I_MPI_FABRICS=tcp
$> mpirun -hostfile $OAR_NODEFILE ...
```

MPI with OAR

- Reservation has to match the required number of MPI processes:
 ↪ **Ex:** `oarsub -l nodes=2/core=12[...]`
- Running MPI code using traditional `mpirun`

```
# Intel MPI
```

```
$> module load toolchain/intel
```

```
# ONLY on moonshot node have no IB card: export I_MPI_FABRICS=tcp
```

```
$> mpirun -hostfile $OAR_NODEFILE ...
```

```
# OpenMPI
```

```
$> module load mpi/OpenMPI
```

```
$> mpirun -hostfile $OAR_NODEFILE -x PATH -x LD_LIBRARY_PATH ...
```

MPI with OAR

- Reservation has to match the required number of MPI processes:
 ↪ **Ex:** `oarsub -l nodes=2/core=12[...]`
- Running MPI code using traditional `mpirun`

```
# Intel MPI
```

```
$> module load toolchain/intel
```

```
# ONLY on moonshot node have no IB card: export I_MPI_FABRICS=tcp
```

```
$> mpirun -hostfile $OAR_NODEFILE ...
```

```
# OpenMPI
```

```
$> module load mpi/OpenMPI
```

```
$> mpirun -hostfile $OAR_NODEFILE -x PATH -x LD_LIBRARY_PATH ...
```

```
# MVAPICH2
```

```
$> module load mpi/MVAPICH2
```

```
$>-launcher ssh -launcher-exec /usr/bin/oarsh -f $OAR_NODEFILE ...
```

MPI launchers (OAR)

[Github](#)

```
$> oarsub -S <scriptname>          # -S: interpret #OAR comments
```

```
#!/bin/bash -l
#OAR -l nodes=2/core=6,walltime=1
if [ -f /etc/profile ]; then
    . /etc/profile
fi
# Load the intel toolchain and whatever MPI module you need
module purge && module load toolchain/intel # or mpi/{OpenMPI/MVAPICH2}
# ONLY on moonshot node that have no IB card: export I_MPI_FABRICS=tcp

### Intel MPI
mpirun -hostfile $OAR_NODEFILE mpi_program

### OpenMPI
mpirun -hostfile $OAR_NODEFILE -x PATH -x LD_LIBRARY_PATH mpi_program

### MVAPICH2
mpirun -launcher ssh -launcher-exec /usr/bin/oarsh \
    -f $OAR_NODEFILE mpi_program
```

MPI with Slurm

sbatch/srun option	Description
-N <N>	number of distributed nodes
--ntasks-per-node=<N>	number of MPI processes per node
-c 1	(default) set a single thread per MPI process
-c <T>	number of OpenMP threads for hybrid runs

- SLURM able to directly launch MPI tasks (`_recommended_`)
 - ↪ ... and initialize of MPI communications
 - ↪ via **Process Management Interface (PMI)** [v2]

```
$> srun -n $SLURM_NTASKS /path/to/mpiprogram
```

MPI launcher (Slurm)

Docs

```
$> sbatch [-p batch] <scriptname>
```

```
#!/bin/bash -l
#SBATCH -N 2                # Use 2 nodes
#SBATCH --ntasks-per-node=28 # Number of MPI process per node
#SBATCH -c 1                # threads per MPI process (1 unless hybrid code)
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

if [ -f /etc/profile ]; then
    . /etc/profile
fi
# Load the intel toolchain and whatever MPI module you need
module purge
module load toolchain/intel # or mpi/{OpenMPI/MVAPICH2}
# export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
srun -n $SLURM_NTASKS /path/to/your/mpi_program
```


Hands-on 2: MPI jobs

Your Turn!

Hands-on 2

ulhpc-tutorials.rtf.d.io/en/latest/parallel/basics/#paralleldistributed-mpi-jobs

- **Reserve** an **interactive** job to launch 6 MPI processes
 - ↪ across two nodes (2x3), for 30 minutes
- Check and **compile** `src/hello_mpi.c`
 - ↪ `bin/{openmpi,intel,mvapich2}_hello_mpi`
- **execute** the generated binaries
- prepare a **launcher script** `runs/launcher.MPI.sh`
- repeat on a more serious program `src/matrix_mult_mpi.c`
 - ↪ `bin/{openmpi,intel,mvapich2}_matrix_mult_mpi`



Summary

- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs**
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)

Hybrid OpenMP+MPI Programs

MPI Suite	module load...	Compilation command (C)
Intel MPI	toolchain/intel	mpiicc -Wall -qopenmp [-xhost] -O2 [...]
OpenMPI	mpi/OpenMPI	mpicc -Wall -fopenmp -O2 [...]
MVAPICH2	mpi/MVAPICH2	mpicc -Wall -fopenmp -O2 [...]

- adapt (and share) OMP_NUM_THREADS environment variable
- **(Slurm only)**: adapt -c <T>: number of OpenMP threads
- **(OAR only)**: you have to take the following elements into account:
 - compute accurately MPI processes per node <PPN>
 - pass it to mpirun:
 - ✓ OpenMPI: mpirun -npnode <PPN> -np <N>
 - ✓ Intel MPI: mpirun -perhost <PPN> -np <N>
 - ✓ MVAPICH2: mpirun -ppn <PPN> -np <N>
 - (**_Intel MPI only_**): I_MPI_PIN_DOMAIN=omp
 - (**_MVAPICH2 only_**) MV2_ENABLE_AFFINITY=0

Hybrid OpenMP+MPI with OAR

Intel MPI

```
mpirun -perhost ${NPERNODE:=1} -np ${NP} \  
-genv OMP_NUM_THREADS=${OMP_NUM_THREADS} \  
-genv I_MPI_PIN_DOMAIN=omp \  
-hostfile $OAR_NODEFILE hybrid_program
```

Hybrid OpenMP+MPI with OAR

Intel MPI

```
mpirun -perhost ${NPERNODE:=1} -np ${NP} \  
-genv OMP_NUM_THREADS=${OMP_NUM_THREADS} \  
-genv I_MPI_PIN_DOMAIN=omp \  
-hostfile $OAR_NODEFILE hybrid_program
```

OpenMPI

```
mpirun -npernode ${NPERNODE:=1} -np ${NP} \  
-x OMP_NUM_THREADS -x PATH -x LD_LIBRARY_PATH \  
-hostfile $OAR_NODEFILE hybrid_program
```

Hybrid OpenMP+MPI with OAR

Intel MPI

```
mpirun -perhost ${NPERNODE:=1} -np ${NP} \  
-genv OMP_NUM_THREADS=${OMP_NUM_THREADS} \  
-genv I_MPI_PIN_DOMAIN=omp \  
-hostfile $OAR_NODEFILE hybrid_program
```

OpenMPI

```
mpirun -npernode ${NPERNODE:=1} -np ${NP} \  
-x OMP_NUM_THREADS -x PATH -x LD_LIBRARY_PATH \  
-hostfile $OAR_NODEFILE hybrid_program
```

MVAPICH2

```
export MV2_ENABLE_AFFINITY=0  
mpirun -ppn ${NPERNODE:=1} -np ${NP} \  
-genv OMP_NUM_THREADS=${OMP_NUM_THREADS} \  
-launcher ssh -launcher-exec /usr/bin/oarsh \  
-f $MACHINEFILE hybrid_program
```

Hybrid OpenMP+MPI with Slurm

```
#!/bin/bash -l
#SBATCH -N 2                # Use 2 nodes
#SBATCH --ntasks-per-node=1 # Number of MPI process per node
#SBATCH -c 4                # Number of OpenMP threads per MPI process
#SBATCH --time=0-01:00:00
#SBATCH -p batch
#SBATCH --qos=qos-batch

if [ -f /etc/profile ]; then
    . /etc/profile
fi
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK}

# Load the intel toolchain and whatever MPI module you need
module purge
module load toolchain/intel    # or mpi/{OpenMPI|MVAPICH2}
# export I_MPI_PMI_LIBRARY=/usr/lib64/libpmi.so
srun -n $SLURM_NTASKS /path/to/your/hybrid_program
```

Hands-on 3: Hybrid OpenMP+MPI

Your Turn!

Hands-on 3

ulhpc-tutorials.rtfld.io/en/latest/parallel/basics/#hybrid-openmpmpi-programs

- **Reserve** an **interactive** job to launch
 - ↪ 2 MPI processes (on 2 nodes) / 4 OpenMP threads
- Check and **compile** `src/hello_hybrid.c`
 - ↪ `bin/{openmpi,intel,mvapich2}_hello_hybrid`
- **execute** the generated binaries
 - ↪ set `$OMP_NUM_THREADS`
 - ↪ compute `$NPERHOST`
- prepare a **launcher script** `runs/launcher.hybrid.sh`



Summary

- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks**
- 6 High-Performance Linpack (HPL)

HPC Interconnect Benchmarking

OSU Micro-Benchmarks Instructions

ulhpc-tutorials.rtfid.io/en/latest/parallel/mpi/OSU_MicroBenchmarks/

- **Pre-requisites:** get an interactive job for compilation

```
### Iris cluster
```

```
(access)$> srun -p interactive -N 2 --ntasks-per-node 1 --pty bash  
# aliased/short version: 'si -N 2 --ntasks-per-node 1'  
# best-effort mode :      append '--qos qos-best-effort'  
# within a reservation: append '--reservation <name>'
```

```
### Gaia, chaos cluster
```

```
(access)$> oarsub -I -l nodes=2/core=1,walltime=4  
# best-effort mode :      append '-t besteffort'  
# within a reservation: append '-t inner=<containerID>'
```

OSU micro-benchmarks

<http://mvapich.cse.ohio-state.edu/benchmarks/>

- We will build **version 5.5 of the OSU micro-benchmarks**
- Focusing on (only) two one-sided benchmarks:
 - ↪ `osu_get_latency` - Latency Test
 - ↪ `osu_get_bw` - Bandwidth Test

Building the Benchmarks

Your Turn!

Hands-on 4

ulhpc-tutorials.rtf.d.io/en/latest/parallel/mpi/OSU_MicroBenchmarks/

- Get the sources
- Uncompress them
- Compilation based on
 - ↪ **Intel MPI**
 - ↪ **Open MPI**
 - ↪ **Open MPI** over Ethernet interface
 - ✓ highlight performance drops compared to Infiniband

Building the Benchmarks

```
$> configure --prefix=<path>; make && make install
```

- Based on Autotools/Automake

- ↪ **Rely on** `--prefix=$(pwd)` to state where to install
- ↪ sometimes being in a separate build directory raise issues!
 - ✓ Ex: *osu_util.h: No such file or directory* (missing header)
 - ✓ then you have to play with `CFLAGS=-I<path>` (or `LDPATH`)

```
$> mkdir <builddir> && cd <builddir>  
$> ../src/configure [CC=<compiler>] --prefix=$(pwd)  
$> make && make install
```

Building the Benchmarks

```
$> configure --prefix=<path>; make && make install
```

- Based on **Autotools**/Automake

- **Rely on** `--prefix=$(pwd)` to state where to install
- sometimes being in a separate build directory raise issues!
 - ✓ Ex: *osu_util.h: No such file or directory* (missing header)
 - ✓ then you have to play with `CFLAGS=-I<path>` (or `LDPATH`)

```
$> mkdir <builddir> && cd <builddir>  
$> ../src/configure [CC=<compiler>] --prefix=$(pwd)  
$> make && make install
```

- Now common to have **CMake** based software

```
$> cmake ../src/  
$> make && make install
```

Running the Benchmarks

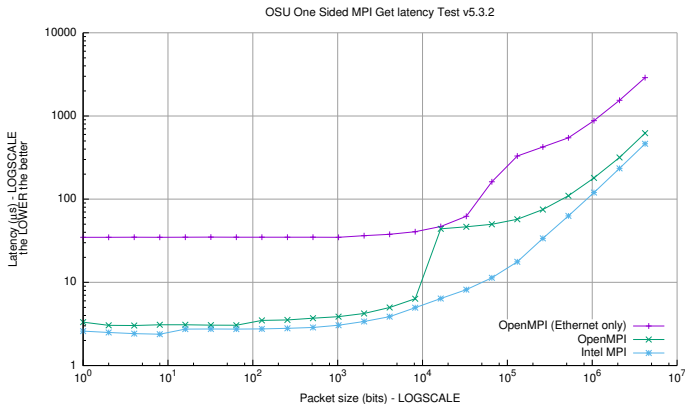
Your Turn!

- Build directory:
libexec/osu-micro-benchmarks/mpi/one-sided/
- Prepare a batch launcher
 - ↪ copy and adapt the **default SLURM launcher**
- Run it in batch mode

```
$> cd ~/tutorials/OSU-MicroBenchmarks/runs
### On iris
$> sbatch ./launcher-OSU.intel.sh osu_get_bw
$> sbatch ./launcher-OSU.intel.sh osu_get_latency
### On gaia, chaos
$> oarsub -S ./launcher-OSU.intel.sh
```

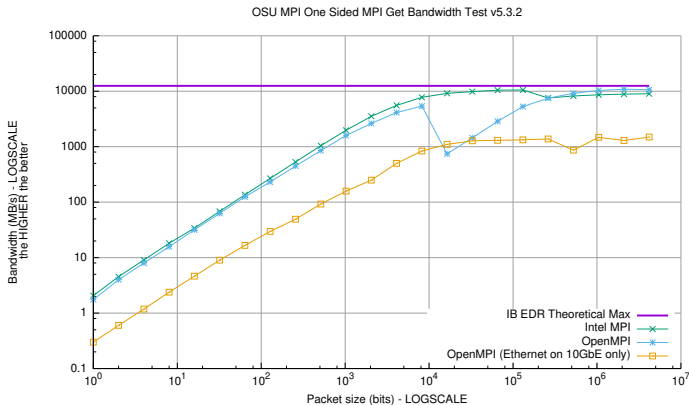
Interconnect Performances

- Based on OSU Micro-benchmarks



Interconnect Performances

- Based on OSU Micro-benchmarks





Summary

- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)**

High-Performance Linpack (HPL)

HPL Instructions

ulhpc-tutorials.rtfid.io/en/latest/parallel/mpi/HPL/

- **Pre-requisites:** get an interactive job for compilation
↪ **Q:** justify the difference with the job request made for OSU.

```
### Iris cluster
(access)$> srun -p interactive -n 14 --pty bash
# aliased/short version: 'si -n 14'
# best-effort mode :    append '--qos qos-best-effort'
# within a reservation: append '--reservation <name>'
```

```
### Gaia, chaos cluster
(access)$> oarsub -I -l nodes=1/core=6,walltime=4
# best-effort mode :    append '-t besteffort'
# within a reservation: append '-t inner=<containerID>'
```



High-Performance Linpack (HPL)

<http://www.netlib.org/benchmark/hpl/>

- Portable implem. of High-Performance Linpack (HPL) Benchmark
 - ↪ for Distributed-Memory Computers
 - ↪ *reference* benchmark for ranking the **Top500** list
- We will build **version 2.2 of the HPL**
 - ↪ Focusing (only) on Intel MPI+MKL build
 - ↪ **Pre-requisites:**
 - ✓ clone **ULHPC/tutorials** and **ULHPC/launcher-scripts** repositories
 - ✓ preparing your working directory

Building HPL

Your Turn!

Hands-on 4

ulhpc-tutorials.rtf.d.io/en/latest/parallel/mpi/HPL/

- Get the sources
- Uncompress them
- Compilation based on the **Intel MPI** suit
 - ↪ Prepare and adapt `src/hpl-2.2/Make.intel64`
- Compile it !

Building HPL

```
$> cd ~/tutorials/HPL/src/hpl-2.2
$> cp setup/Make.Linux_Intel64 Make.intel64
$> vim Make.intel64
# [...] change TOPdir and MP{dir,inc,lib} (at least)
$> make arch=intel64 clean_arch_all
$> make arch=intel64
```

Preparing the HPL Benchmark Run

Your Turn!

- Build directory: `bin/intel64`
- Prepare a batch launcher
 - ↪ copy and adapt the **default SLURM launcher**
- Prepare an input `HPL.dat` file
 - ↪ use **Tuning HPC Online** for some default settings

• Main HPL parameters constraints

- ↪ `PxQ = <nodes>*<cores> = $SLURM_NTASKS`
- ↪ Problem size: `N` (to be as large as possible)
 - ✓ $N = \alpha \sqrt{\#nodes * RAM * 1024}$ where `RAM` is expressed in GiB
- ↪ `NB`: depends on processor architecture (Ex: **Intel MKL notes**)
 - ✓ `NB = 192` on `iris` cluster

Example HPL.dat

```
HPLinpack benchmark input file
Innovative Computing Laboratory, University of Tennessee
HPL.out      output file name (if any)
6            device out (6=stdout,7=stderr,file)
1            # of problems sizes (N)
24650       Ns
1            # of NBs
192         NBs
0           PMAP process mapping (0=Row-,1=Column-major)
2           # of process grids (P x Q)
2 4         Ps
14 7        Qs
[...]
```

- Targeting 1 node in this case on 2 sets of parameters ($P \times Q = 28$)

	N	NB	P	Q
Run 1	24650	192	2	14
Run 2	24650	192	4	7

HPL Benchmark [batch] Runs

- Adapt the default SLURM launcher
- Run it

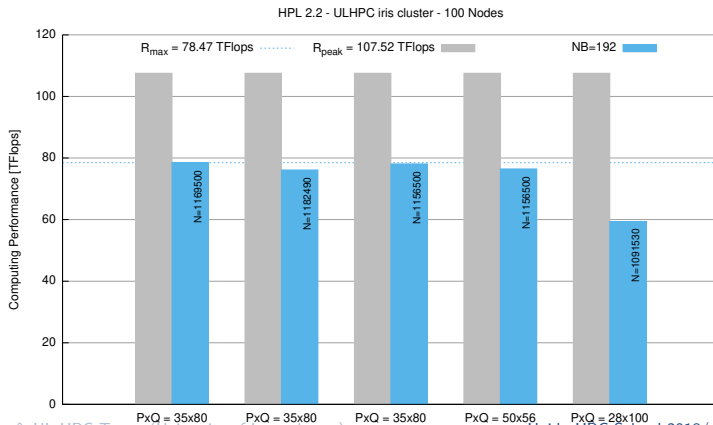
```
$> cd ~/tutorials/HPL/runs
$> cp ../ref.ulhpc.d/HPL.dat .
### On iris
$> sbatch ./launcher-HPL.intel.sh
### On gaia, chaos
$> oarsub -S ./launcher-HPL.intel.sh
```

- Grab the HPL results from the output logs

```
# T/V      N      NB      P      Q      Time      Gflops
$> grep WR slurm-2758.out
WR11C2R4  24650   192     2     14    13.51    7.392e+02
WR11C2R4  24650   192     4      7    12.69    7.869e+02
```

Computing Performances / HPL

- Based on High-Performance Linpack (HPL)
 ↳ reference benchmark for Top 500



Questions?

<http://hpc.uni.lu>

High Performance Computing @ uni.lu

Prof. Pascal Bouvry
Dr. Sebastien Varrette
Valentin Plugaru
Sarah Peter
Hyacinthe Cartiaux
Clement Parisot

University of Luxembourg, Belval Campus
Maison du Nombre, 4th floor
2, avenue de l'Université
L-4365 Esch-sur-Alzette
mail: hpc@uni.lu



- 1 Introduction
- 2 Threaded Parallel OpenMP Jobs
- 3 Parallel/distributed MPI Jobs
- 4 Hybrid OpenMP+MPI Jobs
- 5 OSU Micro-Benchmarks
- 6 High-Performance Linpack (HPL)